

Problem Statement: The data scientists at BigMart have collected sales data for 1559 products across 10 stores in different cities for the year 2013. Now each product has certain attributes that sets it apart from other products.

Breakdown of the Problem Statement: Supervised machine learning problem. The target value will be Item_Outlet_Sales. Aim of the Notebook: The objective is to create a model that can predict the sales per product for each store. Using this model, BigMart will try to understand the properties of products and stores which play a key role in increasing sales

The Big Mart Sales dataset contains information on sales of products in various stores of the Big Mart chain. The dataset contains 12 variables, including both categorical and numerical data. Here's a summary of the variables in the Big Mart Sales dataset

```
import pandas as pd
import numpy as np

filename='/content/Train.csv'
df = pd.read_csv(filename, index_col=0)

df.head()
```

Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility
FDA15	9.30	Low Fat	0.016047
DRC01	5.92	Regular	0.019278
FDN15	17.50	Low Fat	0.016760
FDX07	19.20	Regular	0.000000
NCD19	8.93	Low Fat	0.000000

Item_Identifier	Item_Type	Item_MRP	Outlet_Identifier
FDA15	Dairy	249.8092	OUT049
DRC01	Soft Drinks	48.2692	OUT018
FDN15	Meat	141.6180	OUT049
FDX07	Fruits and Vegetables	182.0950	OUT010
NCD19	Household	53.8614	OUT013

Outlet_Location_Type	Outlet_Establishment_Year	Outlet_Size
Item_Identifier		
FDA15	1999	Medium
Tier 1		
DRC01	2009	Medium
Tier 3		
FDN15	1999	Medium
Tier 1		
FDX07	1998	NaN
Tier 3		
NCD19	1987	High

Tier 3

Item_Identifier	Outlet_Type	Item_Outlet_Sales
FDA15	Supermarket Type1	3735.1380
DRC01	Supermarket Type2	443.4228
FDN15	Supermarket Type1	2097.2700
FDX07	Grocery Store	732.3800
NCD19	Supermarket Type1	994.7052

Lets explore the DataFrame

```
df.shape
```

```
(8523, 11)
```

#finding NaN values

```
df.isnull().sum()
```

```
Item_Weight      1463
Item_Fat_Content      0
Item_Visibility      0
Item_Type          0
Item_MRP           0
Outlet_Identifier      0
Outlet_Establishment_Year  0
Outlet_Size      2410
Outlet_Location_Type      0
Outlet_Type          0
Item_Outlet_Sales      0
dtype: int64
```

so we have some nulls, we should do to deal with them

```
df.describe()
```

```
      Item_Weight  Item_Visibility  Item_MRP
Outlet_Establishment_Year \
count  7060.000000    8523.000000  8523.000000
8523.000000
mean     12.857645         0.066132   140.992782
1997.831867
std       4.643456         0.051598    62.275067
8.371760
min       4.555000         0.000000    31.290000
1985.000000
25%       8.773750         0.026989    93.826500
1987.000000
50%      12.600000         0.053931   143.012800
1999.000000
```

```

75%      16.850000      0.094585      185.643700
2004.000000
max      21.350000      0.328391      266.888400
2009.000000

```

```

      Item_Outlet_Sales
count      8523.000000
mean       2181.288914
std        1706.499616
min         33.290000
25%        834.247400
50%       1794.331000
75%       3101.296400
max       13086.964800

```

```
df.groupby('Item_Type').Item_Weight.agg([min, max, np.mean])
```

	min	max	mean
Item_Type			
Baking Goods	4.880	20.85	12.277108
Breads	4.635	20.85	11.346936
Breakfast	6.425	21.10	12.768202
Canned	4.615	21.35	12.305705
Dairy	4.805	20.70	13.426069
Frozen Foods	4.555	20.85	12.867061
Fruits and Vegetables	5.460	21.35	13.224769
Hard Drinks	4.610	19.70	11.400328
Health and Hygiene	5.175	21.25	13.142314
Household	5.030	21.25	13.384736
Meat	5.150	21.25	12.817344
Others	5.500	20.50	13.853285
Seafood	5.365	20.75	12.552843
Snack Foods	5.095	21.25	12.987880
Soft Drinks	4.590	20.75	11.847460
Starchy Foods	6.695	21.20	13.690731

```
df.groupby(['Item_Type', 'Item_Fat_Content']).mean()
```

```
<ipython-input-192-c5a2ee088b20>:1: FutureWarning: The default value
of numeric_only in DataFrameGroupBy.mean is deprecated. In a future
version, numeric_only will default to False. Either specify
numeric_only or select only columns which should be valid for the
function.
```

```
df.groupby(['Item_Type', 'Item_Fat_Content']).mean()
```

		Item_Weight	Item_Visibility
Item_MRP \			
Item_Type	Item_Fat_Content		
Baking Goods	LF	12.052500	0.068426
115.641430			

121.286145	Low Fat	12.633128	0.066465
133.311137	Regular	11.976126	0.071573
104.131575	low fat	10.060000	0.048614
111.424431	reg	12.825000	0.089004
...	
Starchy Foods LF		14.375000	0.067391
156.902489	Low Fat	13.669692	0.066966
158.575569	Regular	13.757200	0.069425
133.824861	low fat	14.150000	0.037877
124.004600	reg	12.375833	0.062482
147.256086			
Outlet_Establishment_Year			
Item_Outlet_Sales			
Item_Type	Item_Fat_Content		
Baking Goods LF		1995.500000	
1670.791810	Low Fat	1997.491694	
1781.871029	Regular	1998.117647	
2192.281401	low fat	1995.000000	
1345.581800	reg	1999.153846	
1089.504877			
...		...	
Starchy Foods LF		1998.111111	
2466.419111	Low Fat	1998.986111	
2582.980347	Regular	1998.694915	
2141.867315	low fat	1987.000000	
1618.559800	reg	1995.714286	
2177.166000			
[70 rows x 5 columns]			

```

#filling the nan values.
group_means = df.groupby(['Item_Type', 'Item_Fat_Content'])
['Item_Weight'].mean()
df.loc[df['Item_Weight'].isnull(), 'Item_Weight'] =
df.loc[df['Item_Weight'].isnull(), ['Item_Type',
'Item_Fat_Content']].apply(lambda x: group_means[x['Item_Type'],
x['Item_Fat_Content']], axis=1)

df.isnull().sum()

Item_Weight          0
Item_Fat_Content      0
Item_Visibility      0
Item_Type            0
Item_MRP             0
Outlet_Identifier    0
Outlet_Establishment_Year  0
Outlet_Size         2410
Outlet_Location_Type  0
Outlet_Type          0
Item_Outlet_Sales    0
dtype: int64

(df['Item_Visibility']==0).sum()

526

df['Outlet_Size'].value_counts()

Medium    2793
Small     2388
High       932
Name: Outlet_Size, dtype: int64

```

we fill NaN values of Outlet_Size with Mode Imputation

```

df['Outlet_Size'].mode()

0    Medium
Name: Outlet_Size, dtype: object

df['Outlet_Size'].fillna(df['Outlet_Size'].mode()[0], inplace=True)

df['Outlet_Size']

Item_Identifier
FDA15    Medium
DRC01    Medium
FDN15    Medium
FDX07    Medium
NCD19    High

```

```

FDF22      High
FDS36      Medium
NCJ29      Small
FDN46      Medium
DRG01      Small
Name: Outlet_Size, Length: 8523, dtype: object

df['Item_Fat_Content'].value_counts()

Low Fat    5089
Regular    2889
LF          316
reg         117
low fat     112
Name: Item_Fat_Content, dtype: int64

#replacing 2 identical words into one word
df.replace({'LF':'Low Fat', 'reg':'Regular','low fat':'Low Fat'},
inplace=True)

df['Item_Fat_Content'].value_counts()

Low Fat    5517
Regular    3006
Name: Item_Fat_Content, dtype: int64

#use the correlation to find out which column affecting more
df.corrwith(df['Item_Outlet_Sales'])

<ipython-input-203-84f42a0ed7ac>:2: FutureWarning: The default value
of numeric_only in DataFrame.corrwith is deprecated. In a future
version, it will default to False. Select only valid columns or
specify the value of numeric_only to silence this warning.
  df.corrwith(df['Item_Outlet_Sales'])

Item_Weight    0.012432
Item_Visibility -0.128625
Item_MRP        0.567574
Outlet_Establishment_Year -0.049135
Item_Outlet_Sales 1.000000
dtype: float64

```

Now, lets draw some graphs to understand dataframe more clear

```

import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(16,8))
sns.distplot(df['Item_Outlet_Sales'], bins=25)
plt.title("Sales distribution")

```

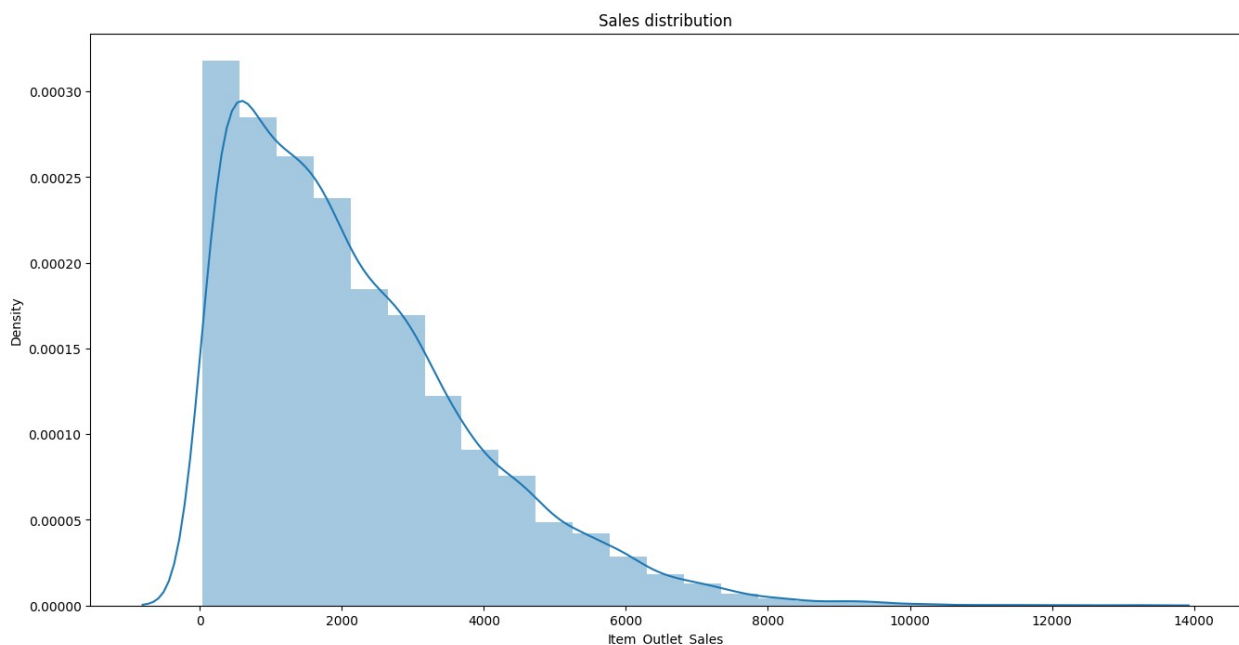
```
<ipython-input-204-187655669e16>:4: UserWarning:
```

```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

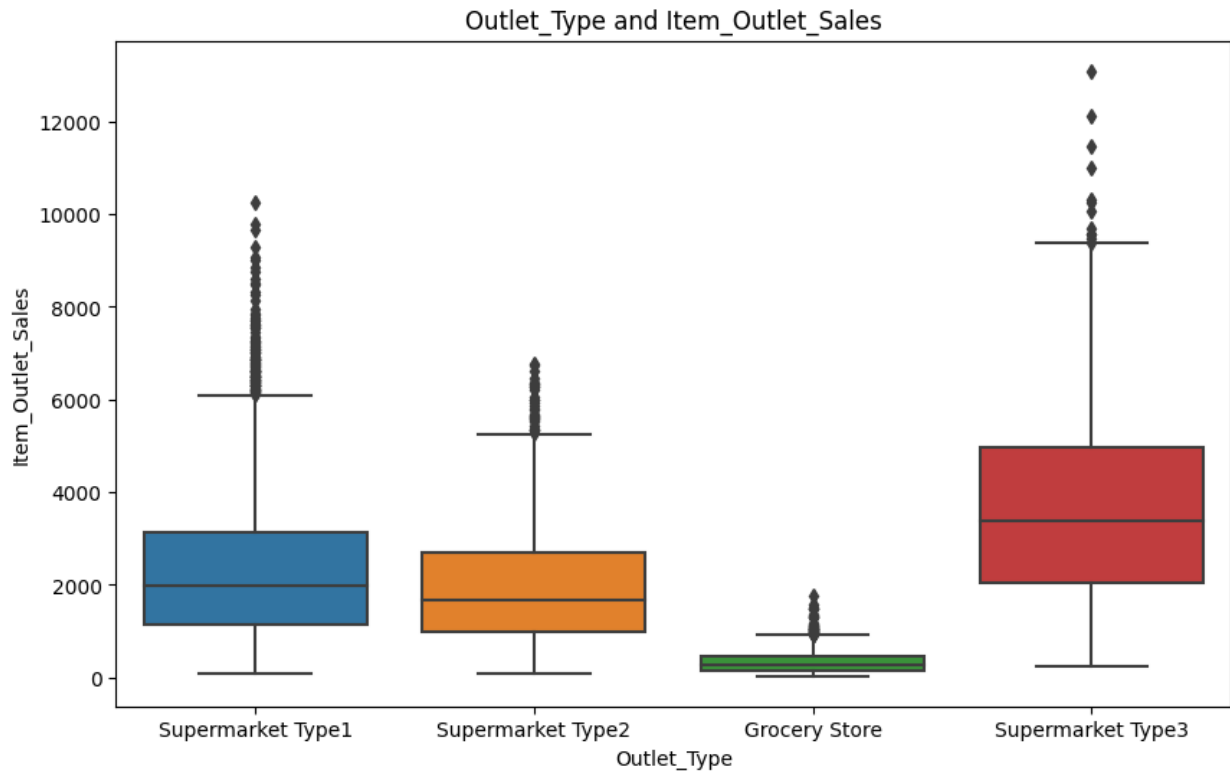
Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df['Item_Outlet_Sales'], bins=25)  
Text(0.5, 1.0, 'Sales distribution')
```



```
plt.figure(figsize=(10,6))  
sns.boxplot(data=df, x='Outlet_Type', y='Item_Outlet_Sales')  
plt.title("Outlet_Type and Item_Outlet_Sales")  
Text(0.5, 1.0, 'Outlet_Type and Item_Outlet_Sales')
```



Using Label Encoder

```
from sklearn.preprocessing import LabelEncoder
labelencoder = LabelEncoder()
columns=['Item_Fat_Content', 'Item_Type', 'Outlet_Identifier', 'Item_Type',
', 'Outlet_Size', 'Outlet_Location_Type', 'Outlet_Type']
for y in columns:
    df[y] = labelencoder.fit_transform(df[y].values)
```

```
df.drop(['Outlet_Identifier'], axis=1)
```

Item_Type \ Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility
FDA15 4	9.300	0	0.016047
DRC01 14	5.920	1	0.019278
FDN15 10	17.500	0	0.016760
FDX07 6	19.200	1	0.000000
NCD19 9	8.930	0	0.000000
...

...			
FDF22	6.865	0	0.056783
13			
FDS36	8.380	1	0.046982
0			
NCJ29	10.600	0	0.035186
8			
FDN46	7.210	1	0.145221
13			
DRG01	14.800	0	0.044878
14			

Item_MRP	Outlet_Establishment_Year	Outlet_Size	\
Item_Identifier			
FDA15	249.8092	1999	1
DRC01	48.2692	2009	1
FDN15	141.6180	1999	1
FDX07	182.0950	1998	1
NCD19	53.8614	1987	0
...
FDF22	214.5218	1987	0
FDS36	108.1570	2002	1
NCJ29	85.1224	2004	2
FDN46	103.1332	2009	1
DRG01	75.4670	1997	2

	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales
Item_Identifier			
FDA15	0	1	3735.1380
DRC01	2	2	443.4228
FDN15	0	1	2097.2700
FDX07	2	0	732.3800
NCD19	2	1	994.7052
...
FDF22	2	1	2778.3834
FDS36	1	1	549.2850
NCJ29	1	1	1193.1136
FDN46	2	2	1845.5976
DRG01	0	1	765.6700

```
[8523 rows x 10 columns]
```

```
X = df.drop('Item_Outlet_Sales', axis=1).values  
y = df['Item_Outlet_Sales']
```

Scaling data

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()  
X = scaler.fit_transform(X)
```

Splitting data

```
#Train/test split  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X,y,  
test_size=0.2, random_state=12)
```

Random Forest Regressor

```
from sklearn.ensemble import RandomForestRegressor
```

```
RF_model = RandomForestRegressor()  
RF_model.fit(X_train, y_train)
```

```
RandomForestRegressor()
```

```
prediction = RF_model.predict(X_train)  
prediction
```

```
array([ 520.768786, 1579.697054, 2437.633618, ..., 2535.273188,  
       1242.502644, 2928.161768])
```

this is final our guess with comparison with the real price

```
pd.DataFrame({'Guess': prediction, 'Real Price':y_train})
```

Item_Identifier	Guess	Real Price
DRM37	520.768786	591.2304
NCQ42	1579.697054	1017.3424
NCE54	2437.633618	2332.2974
NCQ41	3693.771846	3901.5880
FDR19	2627.346670	2770.3938
...
FDK21	507.119886	500.6816
FDS13	4962.713408	5299.7680
FDY19	2535.273188	2828.3184

FDX37	1242.502644	998.7000
NCN26	2928.161768	3866.9664

[6818 rows x 2 columns]

```
from sklearn.metrics import mean_squared_error
mse=mean_squared_error(prediction,y_train)
wrong=np.sqrt(mse)
print(wrong)
```

425.2745791230239

```
RF_accuracy = round(RF_model.score(X_train,y_train) * 100,2)
RF_accuracy
```

93.78

XGBRegressor

```
from xgboost import XGBRegressor

xgb = XGBRegressor(n_estimators=100, random_state=42)
xgb.fit(X_train, y_train)
y_pred = xgb.predict(X_test)
```

```
from sklearn.metrics import mean_absolute_error
mean = mean_absolute_error(y_test,y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
```

```
print(f"RMSE: {rmse}")
print(f'mean: {mean}')
```

RMSE: 1127.9347582596201
mean: 802.7026629318935

#Accuracy

```
accur = round(xgb.score(X_train,y_train) * 100,2)
accur
```

86.67

GradientBoostingRegressor

```
from sklearn.ensemble import GradientBoostingRegressor
gbr = GradientBoostingRegressor()
gbr.fit(X_train, y_train)
gbr_pred = gbr.predict(X_test)
gbr_score = np.sqrt(mean_squared_error(y_test, gbr_pred))

print("RMSE (Gradient Boosting Regressor): ", gbr_score)
```

```
RMSE (Gradient Boosting Regressor): 1056.3519352269534
```

```
#Accuracy
```

```
accur_gbr = round(gbr.score(X_train, y_train)*100,2)  
accur_gbr
```

```
63.03
```

so the final thought of the dataframe, I used Random Forest Regressor(93.7%), XGBRegressor(86.7%), GradientBoostingRegressor (63.3%) accuracy.

Random Forest Regressor is more efficeint among the regressors i used above,

```
import joblib
```

```
# Save the trained model as a .pkl file
```

```
joblib.dump(RF_model, 'model.pkl')
```

```
['model.pkl']
```